

CURRENT AFFAIRS

A data journalism platform for Mexico

Project Vision Document · Efrén Zagal · 2026

1. Vision

Current Affairs will be an independent data journalism platform dedicated to making sense of Mexico through rigorous analysis, transparent methodology, and beautiful interactive storytelling. It sits at the intersection of two major inspirations: the editorial voice and interactive models of FiveThirtyEight, and the infrastructure-first, public-good ethos of Our World in Data.

On the surface, it will be a Mexican FiveThirtyEight — opinionated, readable, data-driven journalism with a clear point of view. Underneath, though, it will be a persistent data platform: live dashboards, self-updating pipelines, and clean datasets that grow more valuable over time. Both layers reinforce each other. An article will say 'as you can see in our spending tracker' and links to the live dashboard. The dashboard links back to the methodology piece. The site has a heartbeat even between publications.

The one-sentence pitch: the intersection of FiveThirtyEight and Our World in Data, for Mexico.

Competitive advantage

Mexican data journalism exists in two buckets. Institutional players — Bloomberg, AS/COA, Wilson Center — do rigorous work but only activate during election cycles, have no independent Mexican voice, and disappear afterward. Independent creators like LaData.mx and El Dato Mx

have the right spirit but lack technical depth and interactive UI. Neither bucket produces the combination of domain expertise, product thinking, and persistent infrastructure that Current Affairs aims for.

The founder brings a genuinely unusual combination: applied mathematics training from ITAM, years of data science and marketplace dynamics work in private industry, and experience leading matching and supply positioning efforts for Uber Latam — all complemented by a Fulbright scholarship to study public policy at Georgetown's McCourt School. The analysis is grounded in real operational expertise. The policy lens is being actively developed. The product instincts come from shipping real systems at scale.

Critically, this is not AI-generated analysis dressed up as journalism. The methodology is transparent, the data is sourced and reproducible, and the interpretations are grounded in domain expertise that cannot be faked. In an environment about to be flooded with superficially credible AI-generated content, that is a genuine differentiator.

2. The Platform

Current Affairs operates as two surfaces simultaneously, each reinforcing the other.

Editorial surface

Articles with embedded charts, narrative-driven, a specific voice and point of view. Written pieces with inline visualizations that support the argument. The goal is accessibility without sacrificing rigor — the same piece should be readable by a curious citizen and credible to a policy expert.

Dashboard surface

Persistent, self-updating views where the data lives. Electoral approval trackers, urban mobility flow maps, public spending explorers. Things people bookmark and return to, not just read once. These are products with users, not charts attached to articles. They get instrumented, iterated, and improved based on what people actually interact with.

The PM-Tech approach

Traditional media outlets approach data journalism as journalists who learned some data tools. Current Affairs approaches it as a product — with infrastructure that scales, pipelines that run without manual intervention, and distribution treated as a growth problem. This mental model is invisible in the output but shows up everywhere in the execution: defensive engineering, iterative UI improvement, thinking about what brings readers back rather than just what gets published.

3. The Six Projects

Six content pillars, each with its own data pipeline, dashboard, and editorial angle. Together they form an ecosystem — each project is standalone, but the data flows between them. A heat wave summer feeds both the climate dashboard and an article on Metro ridership. Economic indicators feed both the spending tracker and the electoral model. Everything talks to everything.

Project 1 — Urban Mobility: Ecobici

STATUS: DATA COLLECTED · PIPELINE NEEDS RESTART · UI TEMPLATE PENDING

The founder collected a proprietary dataset by polling the Lyft/Ecobici API every 8 minutes, running continuously for several weeks, capturing bike availability across all stations in Mexico City. The dataset includes station ID, timestamp, bikes available, and full station metadata including geolocation. This is original data that nobody else has.

The Uber marketplace background applies directly here. Supply positioning, demand forecasting, and equity analysis across a city are problems the founder has solved professionally. The editorial angle is the insider critique — someone who understands how these systems are designed calling out where they fail in practice.

Key article angles

- Equity analysis across alcaldías — does Ecobici serve all of the city equally, or is it concentrated in wealthier areas like Roma, Condesa, and Polanco?
- Flow patterns — how does the city move throughout the day, week, and across seasons?
- Supply gaps — where are bikes missing when people need them most? Using hex-based analysis similar to Uber's internal tooling.

- Forecasting availability — can we predict station-level availability? Useful for riders and operationally interesting.

Strategic role: launch project. Requires the least additional data engineering, lets the full pipeline — scraping, cleaning, analysis, interactive maps, article, publishing, distribution — be tested end to end before stakes are high. The Ecobici dataset also serves as the UI sandbox for building reusable chart, map, and table templates.

Project 2 — Urban Mobility: Metro CDMX

STATUS: CONCEPT · DATA PIPELINE TO BE DESIGNED

A natural sequel to Ecobici, the Metro is a much larger and more politically charged system. Chronic underfunding, overcrowding on key lines, equity across the city, and the 2021 Line 12 collapse that killed 26 people — the Metro sits at the intersection of urban policy, infrastructure, and public accountability in ways Ecobici does not.

CDMX government publishes open data on Metro operations through its open data portal. The pipeline is more tractable than expected. Combined with the reusable UI components built during the Ecobici project, this can move faster.

Key article angles

- Line-by-line performance — on-time rates, delay patterns, capacity utilization
- Equity across the city — which lines serve underserved communities, and how does service quality compare?
- Budget vs. service — tracking infrastructure spending against service quality outcomes
- The Line 12 story — a data-driven accountability piece on what the numbers said before and after the collapse

Together, Ecobici and Metro form a coherent urban mobility beat that nobody in Mexican data journalism currently owns. This is the territory to claim.

Project 3 — Electoral Model

STATUS: CONCEPT · RESEARCH PHASE · SLOW BURN FLAGSHIP

The marquee project and the one most visible to a general audience. Mexico just came off the 2024 election cycle with Claudia Sheinbaum winning the presidency with 59% of the vote — outperforming most polls by several points. The polling credibility gap is real and well-documented, and it is the opening.

The technical foundation is strong: bias correction, sampling strategies, confidence intervals, Monte Carlo simulations, and experience with competitive intelligence models from Uber. The founder has previously built an electoral analysis — modeling how Mexican elections would behave under an electoral college system — that received significant media coverage.

Signal architecture

- Structural layer: economic fundamentals — inflation, USD/MXN exchange rate, Banxico reference rate, unemployment
- Leading indicator: presidential and party approval ratings, tracked continuously
- High-frequency signal: media sentiment scraped from major outlets — Animal Político, Milenio, El Universal, Reforma, El País
- Electoral history: state-by-state composition going back multiple cycles

Key article angles

- Polling accuracy audit — how well have Mexican polls performed in 2006, 2012, 2018, 2024? Who got it right and who did not?
- The economy and the vote — does inflation predict approval? Does approval predict vote share?
- State-by-state electoral composition — which states are actually competitive and which are structural locks?
- Methodology transparency — publishing the model architecture openly, as 538 did

Strategic role: slow burn flagship. Needs groundwork, data sourcing, and credibility-building before the live model launches. Start research and data collection early, publish methodology pieces as the model develops, launch the interactive tracker when the 2027 midterm cycle begins heating up.

Project 4 — Public Spending

STATUS: CONCEPT · DATA ENGINEERING CHALLENGE · LONG-TERM INFRASTRUCTURE

The most civic-minded project and the one that compounds most over time. Government spending data technically exists — in SHCP budget documents, INEGI reports, and Cuenta Pública filings — but it is dense, fragmented, and in formats that make it practically inaccessible to most people. The project's value proposition is exactly what Our World in Data does: take data that exists but is inaccessible and make it legible, trustworthy, and beautiful.

Every year the database grows more valuable. By year five it is a multi-year longitudinal dataset that the founder built and curates — that is infrastructure, not just a project.

Key article angles

- Cross-administration comparison — how has the composition of spending actually changed from PRI to PAN to PRD to Morena? Education vs. security vs. social programs vs. debt service
- Rhetoric vs. reality — does each party spend on what it says it prioritizes?
- The discretionary layer — tracking spending that is hard to see in official summaries
- Morena's fiscal footprint — a detailed analysis of budget priorities under AMLO and Sheinbaum

Strategic role: begin data engineering early — even during the Uber weeks — because cleaning Hacienda data will take time. Publish third, after McCourt has provided more public policy vocabulary. The longer the wait, the better the analysis.

Project 5 — Climate Trends

STATUS: CONCEPT · PURELY DESCRIPTIVE · HIGH ACCESSIBILITY

A descriptive analysis project tracking historical weather data for Mexico — principally Mexico City — to answer a question millions of people are asking every summer: is it actually getting hotter, or does it just feel that way? The answer, backed by clean data and good visualizations, is genuinely useful and deeply relevant.

This is the least politically sensitive project, making it good for broadening the audience beyond policy and politics readers. The Our World in Data comparison is most apt here — honest, well-sourced descriptive analysis without overreach into climate modeling.

Key article angles

- Hottest summers on record — a clean ranking with full historical context
- Rainfall trends — is CDMX getting drier? When are the rains coming compared to historical norms?
- The urban heat island — how does temperature vary across different parts of the city?
- Decade-by-decade shifts — visualizing long-term climate change at the city level

The pipelines feed the editorial work naturally. Every record-breaking heat wave generates automatic context from the dashboard. Every rainy season anomaly is already in the data.

Project 6 — Mexico City Marathon

STATUS: DATA AVAILABLE · ONE-TIME SCRAPING · HIGH VIRAL POTENTIAL

The fun one, and strategically underrated. Race results for the Mexico City Marathon are published publicly by year. A scraping and cleaning project turns them into a longitudinal dataset covering thousands of runners across multiple years. The analysis is purely descriptive but the subject matter has enormous built-in audience — Mexico City's running community is large, active, and highly engaged on social media.

This project diversifies the platform's audience deliberately. Someone who finds Current Affairs through the marathon piece discovers the electoral model and the spending tracker. It also gives the site personality — it is not only a serious policy platform, it is a place that cares about how the city lives, moves, and runs.

Key article angles

- The wall at kilometer 32 — does the data show pace collapse at the expected point in the race?
- Is Mexico City running faster? — finishing time trends across years
- Who finishes and who does not? — DNF patterns by age group, gender, and weather conditions on race day
- Mexican vs. international runners — a comparative profile

4. Roles & Gaps

Building Current Affairs requires five distinct roles. Some are covered, some are gaps, and some will be filled over time.

Data scientist — Covered

The core analytical work: modeling, statistical inference, bias correction, forecasting, competitive intelligence frameworks. This is the founder's primary expertise and the backbone of the platform's credibility.

Data engineer — In progress

Building and maintaining data pipelines that run without manual intervention. Scheduled scripts, API polling, web scraping, and light orchestration via GitHub Actions or cron jobs. The key design question is which data requires continuous pipelines versus one-off collection. This role needs to be built out deliberately during the pre-Georgetown window, with a focus on Python automation and pipeline reliability.

UI engineer / designer — Gap, AI-assisted

The most significant skills gap. The founder works primarily in R, which has a ceiling aesthetically and interactively. The target stack is React-based with D3, Plotly, or Mapbox for visualizations, and a Next.js site that handles both article pages and dashboards. The strategy is not to become a frontend engineer but to develop enough fluency to direct AI tools effectively — using the Ecobici dataset as a sandbox to build reusable component templates before the other projects are ready.

Media & distribution — To be built

Understanding what other outlets are covering, identifying gaps, mapping the information ecosystem, and building an audience. This is not influencer work — it is analytical work applied to the media landscape itself. The minimum viable presence is X/Twitter (where Mexican journalists, policy people, and data practitioners live) and LinkedIn. Distribution should be treated as a growth problem: funnels, retention, what brings people back. The goal is not viral posts but a growing community of people who trust the work.

Public policy — Developing

The McCourt School at Georgetown provides the formal frameworks, vocabulary, and peer network to develop genuine policy fluency. The founder already has strong intuitions from operational work — understanding resource allocation, equity in service distribution, and the gap between policy intent and operational reality. The master's program translates those intuitions into academic frameworks. This role will be significantly stronger by the time the Electoral Model and Public Spending projects are at publication stage.

5. Competitive Landscape

The Mexican data journalism and policy analysis space has three categories of players, each with clear limitations that Current Affairs is positioned to address.

Institutional players

Bloomberg, AS/COA, Wilson Center, and Expansión Política produce rigorous work during election cycles — Bloomberg's 2024 tracker used weighted local regression with pollster accuracy indexes going back to 2006. But these are foreign or institutional voices that activate episodically and disappear. There is no independent Mexican platform doing this work year-round.

Independent Mexican creators

LaData.mx was founded by physicists who wanted to bring data visualization to Mexican journalism and demonstrate it could be done well. Their ambition was to publish daily data stories and educate the public on reading data. El Dato Mx is a solo creator doing analysis on sports and politics with some map visualizations. Both have the right spirit. Neither has the technical depth, interactive UI, or persistent infrastructure that Current Affairs aims for. LaData.mx appears to have slowed significantly.

Government platforms

DataMéxico is the Secretaría de Economía's platform for exploring economic, social, and occupational data through interactive visualizations. Useful as a data source. Zero editorial voice. Not a competitor.

The gap

Nobody is doing what Current Affairs describes: technically rigorous, independently voiced, beautifully interactive, always-on coverage of mobility, public finance, electoral dynamics, climate, and urban life simultaneously. That space is open.

6. Inspirations

FiveThirtyEight (original)

The editorial model. Interactive trackers that people return to daily, not just read once. Methodology transparency as a core product feature. An accessible but rigorous voice that treats readers as intelligent adults. The model pages — where you could flip a state and watch win probability update — were the thing that made it addictive.

Our World in Data

The infrastructure model. Data that technically exists made legible, trustworthy, and beautiful. Honest about uncertainty and limitations. A public good that grows more valuable every year. Well-sourced, reproducible, and built to last.

The Pudding

Visual essays with original data collection and highly unconventional formats. Proof that serious data work can also be playful and surprising.

El País Datos

The best example of Spanish-language data journalism done with genuine production quality. A reference for tone, visual style, and what a serious bilingual data publication looks like.

Bloomberg Graphics

The gold standard for interactive financial and political data visualization. The 2024 Mexico election tracker is the closest existing product to the electoral model Current Affairs wants to build — and it was built by a foreign outlet with a team of people.

7. Timeline & Milestones

Phase 1 — Foundation (Weeks 1–6, while at Uber)

Evenings and weekends only. Focus on laying groundwork, not shipping.

- Complete this project vision document and prepare Google Slides deck
- Map competitive landscape and identify key outlets to monitor
- Explore data availability: INEGI, INE, Hacienda, CDMX open data portal
- Restart Ecobici API polling with a cleaner, more intentional pipeline
- Begin learning Python job automation and GitHub Actions for scheduled tasks
- Identify domain name candidates and set up social media handles

Phase 2 — Building (Weeks 7–14, pre-Georgetown)

The primary building window. Eight weeks of focused work with no job and lighter logistics obligations.

- Ecobici: complete analysis, build interactive map and chart components, write first article, publish
- UI template: reusable components for maps, charts, and tables built on Ecobici data
- Marathon: scrape and clean historical results, publish second piece
- Electoral: begin data sourcing, pollster accuracy audit, methodology research
- Public spending: begin Hacienda data ingestion and cleaning
- Site: live in some form before Georgetown begins

Phase 3 — Sustaining (Georgetown year and beyond)

Pipelines run continuously. Publishing cadence is realistic — one piece per month is more valuable than an ambitious schedule that collapses. Metro project begins. Electoral model builds

toward midterm cycle. Public spending tracker grows with each budget cycle. Climate data updates automatically. The site has a heartbeat independent of publication frequency.

8. Long-term Vision & Business Model

The floor (5 years)

A well-maintained data platform with six live dashboards, a body of rigorous published analysis, and a demonstrable track record as an independent voice on Mexican public life. Enough to meaningfully differentiate in any job application — government, think tanks, private sector, international organizations. A genuine career asset that nobody else in the field has.

The ceiling (5 years)

The definitive data reference for Mexican public life. The platform journalists cite when they need the clean number. The site policy people share internally. A recognized independent voice that international outlets seek for the Mexican data perspective. Potentially a small media business with institutional partnerships, consulting revenue, and one or two collaborators.

The consulting model

The free, open work is the marketing. Every dashboard, every article, every pipeline made public is a demonstration of capability that no CV can replicate. The consulting client does not hire because of a proposal — they hire because they have been reading the site for six months and this is already the most credible voice they know on Mexican mobility data or electoral dynamics.

Government agencies working on urban mobility see the Ecobici and Metro analysis and want this done internally. Political campaigns see the electoral model and want something proprietary. International organizations see the spending tracker and want a custom version. The consulting that flows from this kind of platform is higher quality work with better clients than cold business development ever generates.

The rule that makes it work: the analysis is always independent, the skills are for hire. The moment the site feels editorially for sale, both the open work and the consulting value collapse together.

9. Name & Identity

Current Affairs is a working alias, named after a track from Lorde's album *Virgin*. The name is intentionally double-meaning: current as in contemporary, timely, and relevant — and current as in flow, movement, the thing that runs through a city's veins. It fits the urban mobility work, the electoral tracking, the climate data. It has the right register: serious but not stuffy, international but with a specific place in mind.

A permanent name, visual identity, and domain should be decided before launch but after the content direction is clearer. The identity will emerge from what gets published and how it sounds, not from a branding exercise done in a vacuum. The one early decision worth making: own the X/Twitter handle and domain before someone else does.

Bilingual by design. Spanish for Mexican relevance and domestic job market positioning. English for international reach and Georgetown-era writing practice. The two languages are not translations of each other — they may develop distinct voices for distinct audiences.

End of document.

This is a living document. Update it as the project evolves.